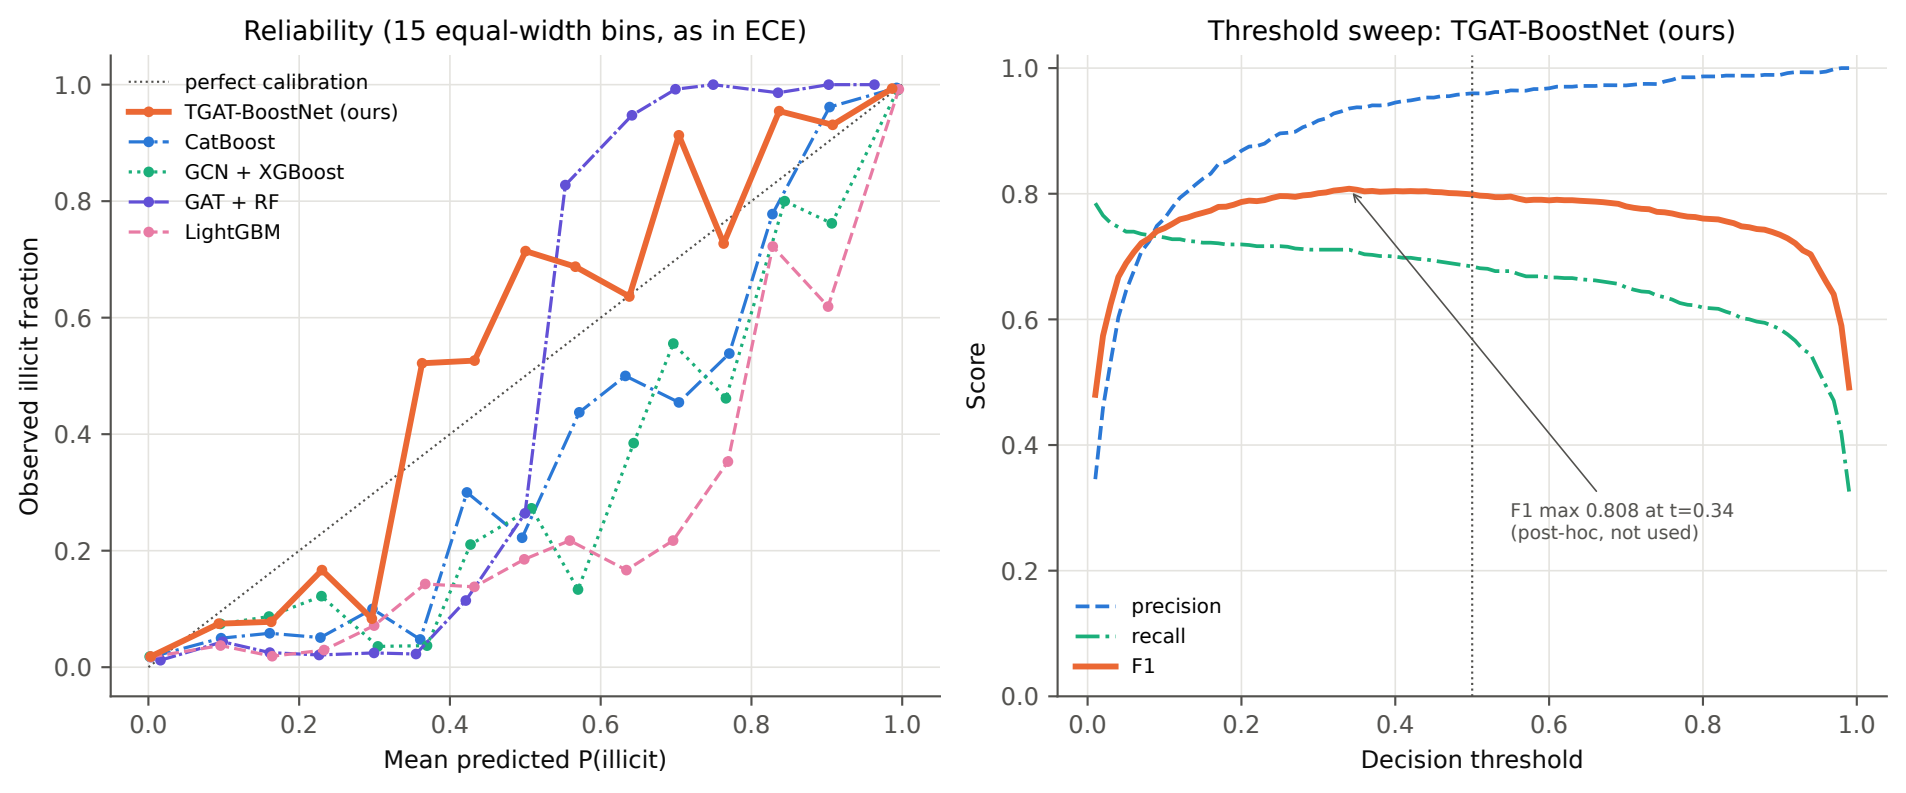
Reliability (15 equal-width bins, as in ECE)
perfect calibration
TGAT-BoostNet (ours)
CatBoost
GCN + XGBoost
GAT + RF
LightGBM
Observed illicit fraction
Mean predicted P(illicit)
0.0
0.2
0.4
0.6
0.8
1.0
Threshold sweep: TGAT-BoostNet (ours)
Score
Decision threshold
F1 max 0.808 at t=0.34
(post-hoc, not used)
precision
recall
F1